

Microcontroller programming

Fuszenecker Róbert
2011.

Embedded systems

An embedded system is a

- computer system
- designed to perform one or few dedicated functions
- often with real-time computing constraints.

Definition: from Wikipedia.

Embedded systems

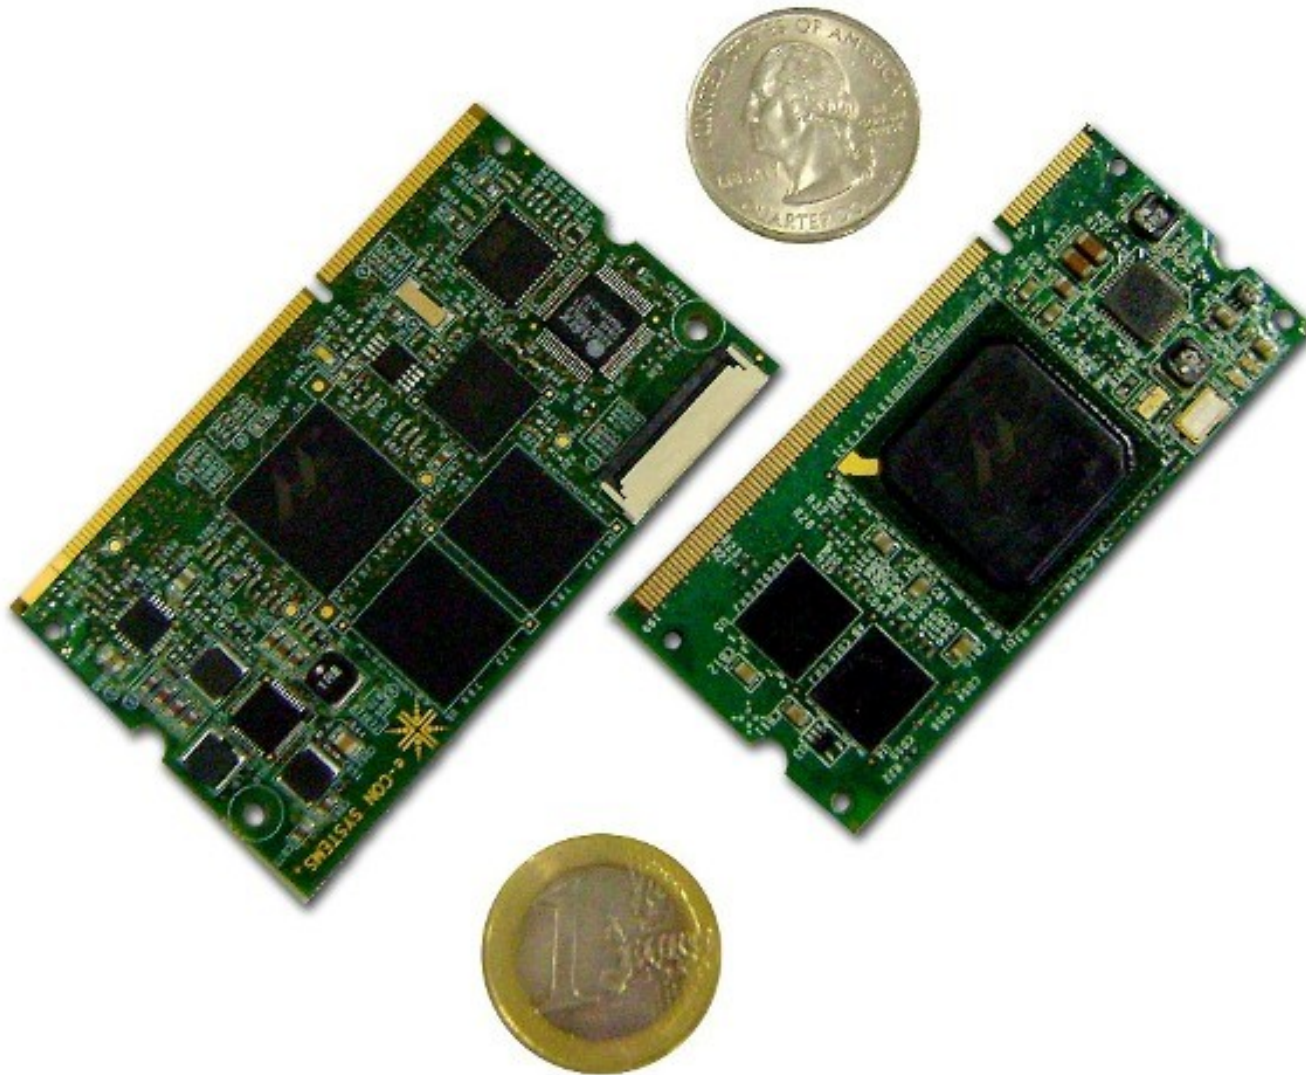
Typical types of embedded systems:

- Industrial computer,
- System on module, computer on module,
- Microcontroller,
- FPGA.

Industrial computers





System on modules



Microcontrollers



FPGAs



FPGA: Field Programmable Gate Array

Embedded operating systems

- Industrial computers and system-on-modules are based on 32-bit processors (with MMU) and have enough resources to run operating systems, i.e. Linux or Windows CE.
- Microcontrollers sometimes run real-time operating systems: FreeRTOS, QNX Neutrino, etc.
- FPGAs generally don't run anything, in certain cases they realize soft microprocessors that may run Linux: MicroBlaze, NIOS 2.

ARM cores

ARM Holdings don't produce any silicon chips, they design only the core of the processors.




ARM: Advanced RISC Machine.

ARM cores


There are 3 kinds of ARM cores:

- Application processors (can run Linux).
 - Cortex-A8 (single core), Cortex-A9 (dual or quad core),
 - Cortex-A15 (quad core, servers),
- Microcontroller cores.
 - Cortex-M3, Cortex-M1 (low power), Cortex-M0 (FPGA),
- Obsolete cores (not recommended for new designs),
 - ARM7, ARM9, ARM11.

ARM cores (Cortex-A8)



ARM cores (Cortex-M3)



Processor modes (Cortex-A)

- Abort Mode,
- Fast Interrupt Mode (FIQ),
- Interrupt Request Mode (IRQ),
- Supervisor Mode (kernel),
- System Mode (similar to User Mode),
- Undefined Mode,
- User Mode (user space applications),
- Monitor Mode (debug).

Processor modes (Cortex-M)

- Thread mode (to execute application software),
- Handler (interrupt and exception handlers),
- Supervisor mode (after RESET).

Generally, “thread mode” is not used in embedded systems.

Register set

ARM processors have 16×32 bit registers, many of them has special functions:

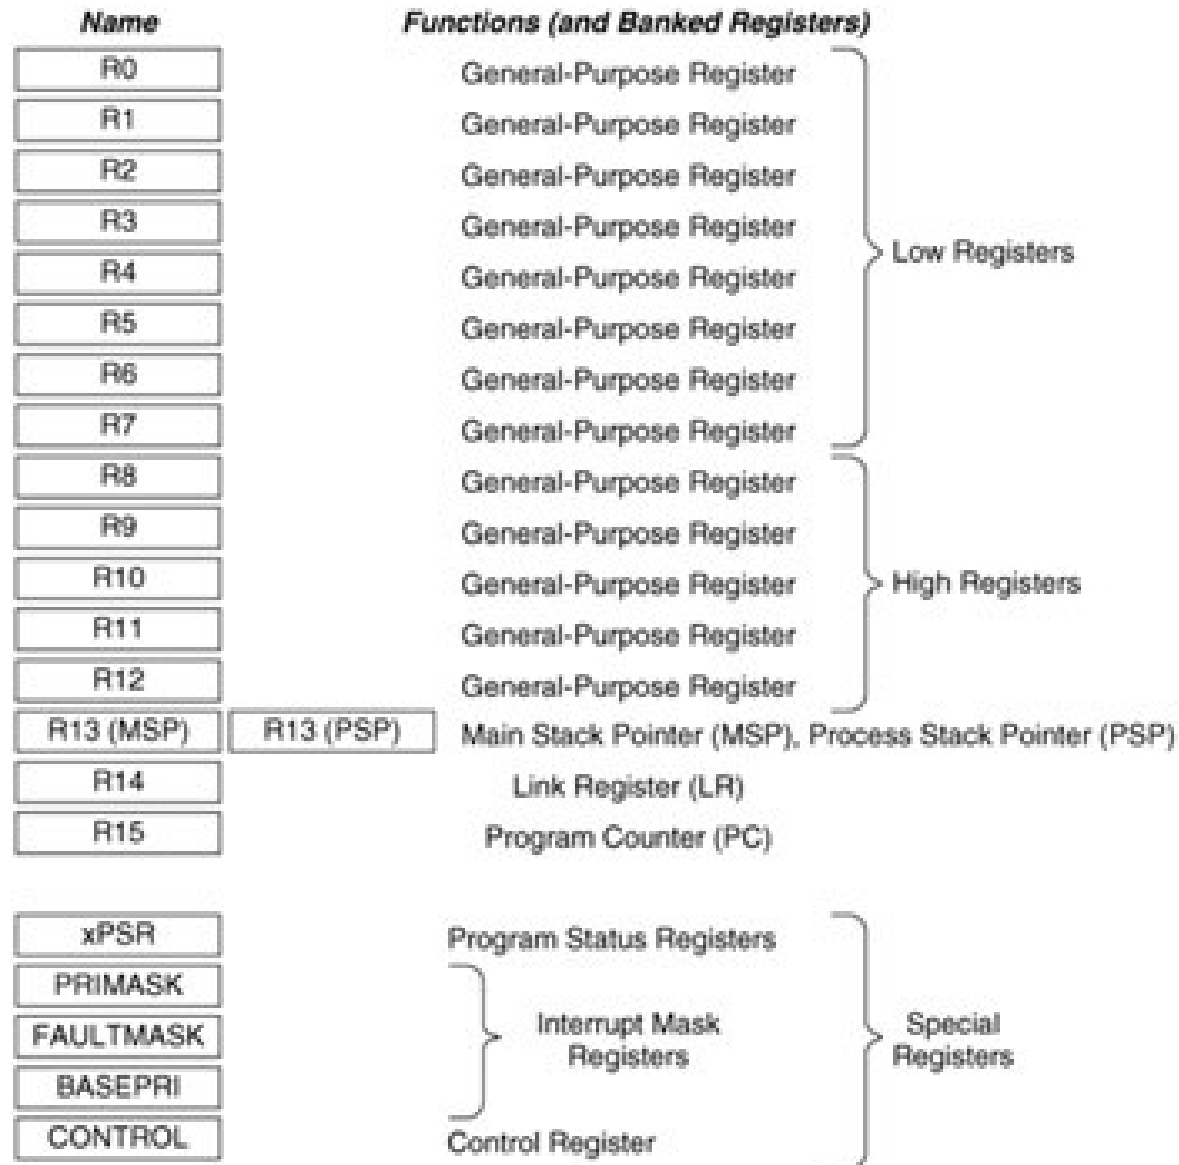
- R15: instruction pointer,
- R14: link register,
- R13: stack pointer,
- R11: frame pointer (optional).

The other ones are general purpose registers.

Register set (Cortex-A)

System & User	FIQ	Supervisor	Abort	IRQ	Undefined
R0	R0	R0	R0	R0	R0
R1	R1	R1	R1	R1	R1
R2	R2	R2	R2	R2	R2
R3	R3	R3	R3	R3	R3
R4	R4	R4	R4	R4	R4
R5	R5	R5	R5	R5	R5
R6	R6	R6	R6	R6	R6
R7	R7	R7	R7	R7	R7
R8	R8_fiq	R8	R8	R8	R8
R9	R9_fiq	R9	R9	R9	R9
R10	R10_fiq	R10	R10	R10	R10
R11	R11_fiq	R11	R11	R11	R11
R12	R12_fiq	R12	R12	R12	R12
R13	R13_fiq	R13_svc	R13_abt	R13_irq	R13_und
R14	R14_fiq	R14_svc	R14_abt	R14_irq	R14_und
R15 (PC)	R15 (PC)	R15 (PC)	R15 (PC)	R15 (PC)	R15 (PC)

Register set (Cortex-M)



Instruction set

- **ARM Cortex-A** cores support both ARM (32 bit) and THUMB (16 bit) instructions (as well as many floating point, DSP and SIMD instructions).
- **ARM Cortex-M** cores support THUMB2 (16 or 32 bit) instructions only.

THUMB2 instructions are the 16 or 32 bit equivalent of the 32-bit ARM instructions. They are used to save memory in resource-critical applications (μ C).

Instruction set

- ARM instruction example (always encoded as a 32-bit word):

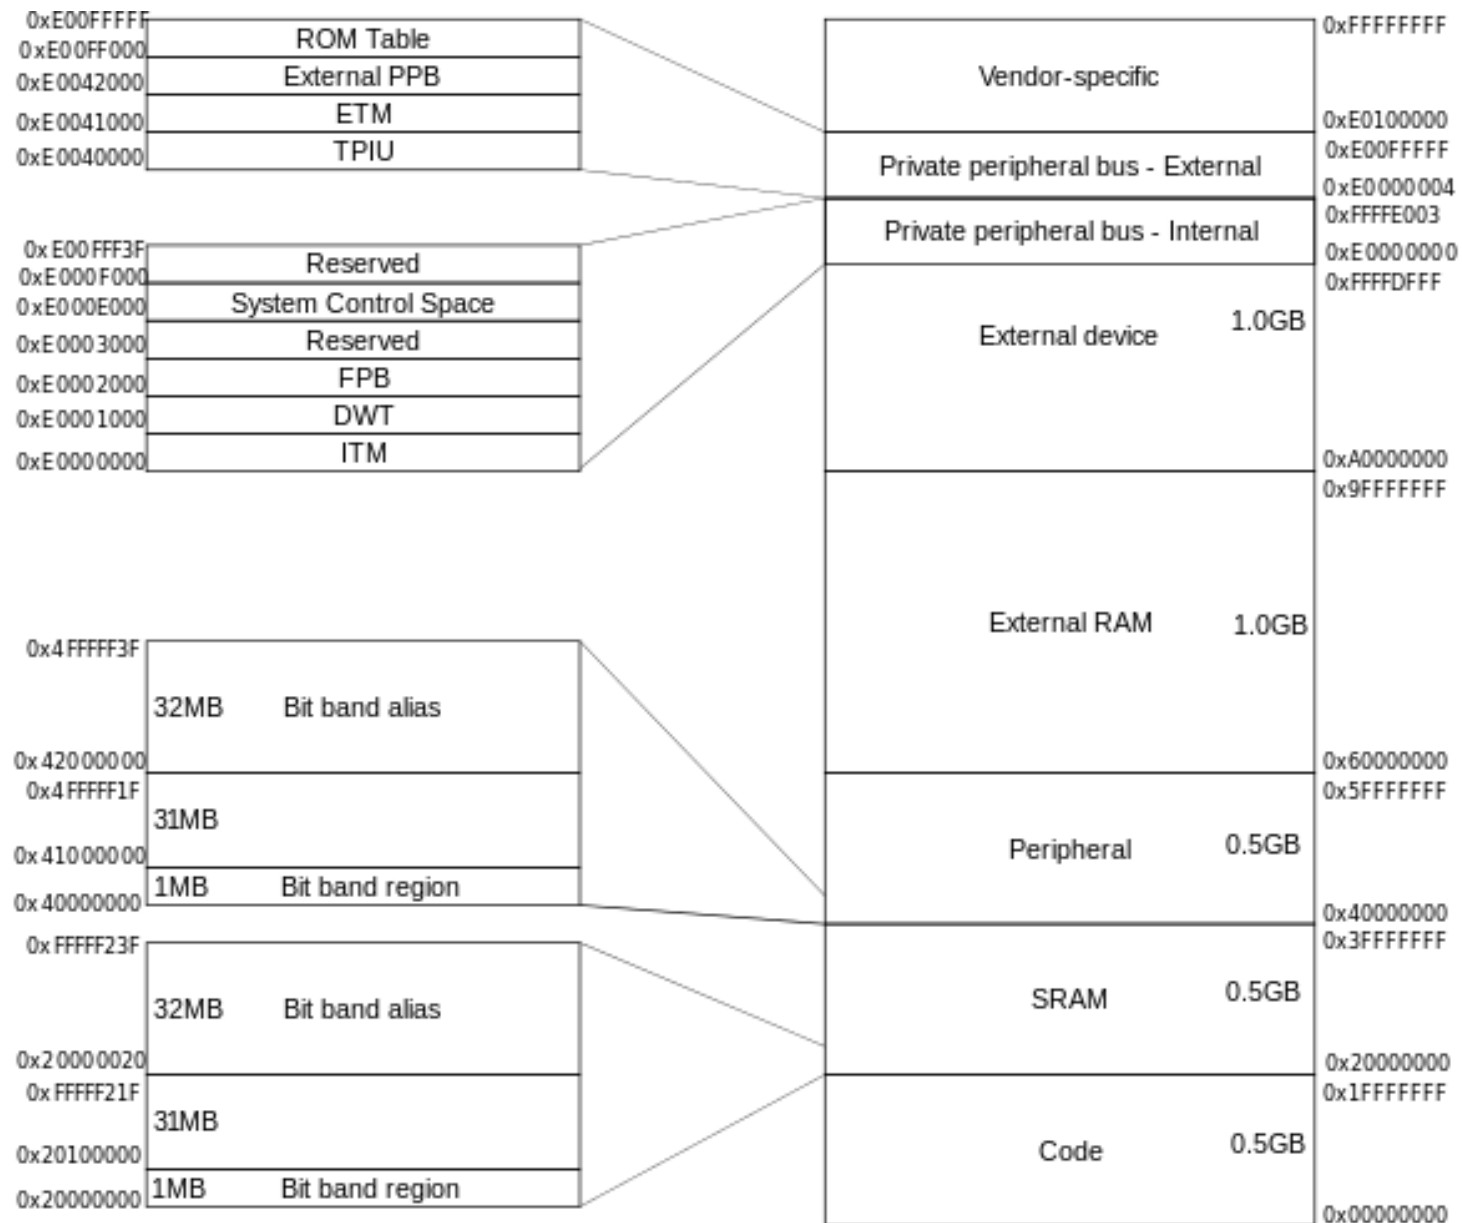
e08631a7 add r3, r6, r7, lsr #3

- THUMB2 instruction examples (16 or 32 bit, depends on the instruction itself):

683a ldr r2, [r7, #10]

fbcb 1204 smlal r1, r2, r3, r4

Memory map (Cortex-M)



Interrupt vectors (Cortex-A)

Physical Address	Vector name
0x0000'0000	Reset
0x0000'0004	Undefined instruction
0x0000'0008	Software interrupt
0x0000'000C	Prefetch abort
0x0000'0010	Data abort
0x0000'0014	(not used)
0x0000'0018	IRQ
0x0000'001C	FIQ

Interrupt vectors (Cortex-M)

Physical Address	Vector name
0x0000'0000	Initial SP
0x0000'0004	Reset Handler
0x0000'0008	NMI Handler
0x0000'000C	Hard Fault Handler
0x0000'0010	Memory Management Handler
0x0000'0014	Bus Fault Handler
0x0000'0018	Usage Fault Handler
...	...
0x0000'0040	Hardware Interrupt 0
0x0000'0044	Hardware Interrupt 1
...	...
0x000'003FC	Hardware Interrupt 243

Booting the processor (Cortex-A)

- After the RESET event the **Cortex-A*** processor starts to execute the **instruction** that can be found at address 0x0000'0000.
- This is typically an unconditional jump to the ROM code that loads the “real” boot loader.

00000000 ldr pc, [pc, #ROM_CODE]

- The boot loader should be able to load and start the Linux kernel or anything else.

Booting the processor (Cortex-M)

- The 0th element of the vector table is the initial value of the stack pointer: usually the last addressable cell of the internal RAM.
- The 1st item is the RESET vector, this is an **ADDRESS** that points to the “main” function.

Booting the processor (Cortex-M)

```
__attribute__((section(".interrupt_vector")))
void (*const Vectors[]) (void) = {
    (void *)0x20003FFC // The initial stack pointer
    Reset_Handler,    // The reset handler
    ...               // Other interrupt handlers
};

extern int main(void);

void Reset_Handler(void) {
    While(1) {
        main();
    }
}
```

The main() function

- Since the stack pointer is set up by the hardware, the processor core is ready to start the “real” task.
- Typically the main() function
 - sets up the hardware peripherals (by using the Firmware library),
 - the interrupt controller,
 - optionally the Memory Protection Unit,
 - and waits for interrupts (low power mode).

Memory mapped I/O

- Several “memory cells” are directly connected to hardware peripherals.
- Hardware registers can be accessed by (memory) read and write instructions, similarly to RAM operations.
- Hardware registers can be protected by the Memory Management Unit (permission control) or the Memory Protection Unit (MPU).

Accessing HW registers

```
#define ADC_BASE_ADDRESS      0x4000A000
#define ADC_START_CONVERSION 0x00000004

typedef struct {
    uint32_t  adc_status;
    uint32_t  adc_value;
    uint32_t  adc_error;
} adc_hw_t;

adc_hw_t *adc0 = (adc_hw_t *)ADC_BASE_ADDRESS;
adc0->adc_status |= ADC_START_CONVERSION;
return_value = adc0->adc_value;
```

Firmware library

- The firmware library is provided by the chip manufacturer (hardware-dependent).
- Contains high level interface to hardware registers:
 - Names instead of addresses (#define),
 - Functions to manipulate the HW registers.
- Sometimes a digital signal processing library is included, so that the programmer doesn't need to work with assembly instructions.

Firmware library example

```
int main()  
{  
    SysCtlPeripheralEnable (SYSCTL_PERIPH_GPIOA);  
  
    GPIOPadConfigSet (GPIO_PORTA_BASE, GPIO_PIN_3,  
        GPIO_STRENGTH_8MA, GPIO_PIN_TYPE_OD);  
  
    GPIOPinTypeGPIOOutput (GPIO_PORTA_BASE, GPIO_PIN_3);  
}
```

Firmware library example

```
SysTickPeriodSet (SYS_TICK_PERIOD);  
SysTickIntRegister (systick_handler);  
SysTickIntEnable();  
SysTickEnable();  
IntMasterEnable();  
  
for (c = 0; ; c++) {  
    counter += c;  
    asm("wfi");  
}  
  
}
```

Firmware library example

```
volatile unsigned int pin_data = 0;
```

```
volatile int counter;
```

```
void systick_handler()
```

```
{
```

```
    pin_data = GPIO_PIN_3 - pin_data;
```

```
    GPIOPinWrite(GPIO_PORTA_BASE, GPIO_PIN_3, pin_data);
```

```
    counter++;
```

```
}
```

With “volatile”

```
for (c = 0;; c++) {
```

```
    counter += c;
```

```
    104: 681a          ldr r2, [r3, #0]
```

```
    106: 18a2          adds r2, r4, r2
```

```
    108: 601a          str r2, [r3, #0]
```

```
asm("wfi");
```

```
    10a: bf30          wfi
```

```
    10c: f104 0401        add.w r4, r4, #1
```

```
    110: e7f8          b.n 104 <main+0x64>
```

R2: counter

R3: &counter

R4: c

Without “volatile”

```
for (c = 0;; c++) {  
    counter += c;  
    asm("wfi");  
    fc: bf30          wfi  
    fe: e7fd          b.n fc <main+0x5c>
```

Compiling the sources

- Commercial compilers:
 - ARM RVDS,
 - Keil,
 - CrossWorks,
 - etc.
- Open source compilers:
 - gcc,
 - clang/llvm.

Compiling the sources (gcc)

With cross-compiled gcc:

- Compiling the source code:

```
gcc -march=arm -c -o crt.o crt.c
```

- Linking the objects:

```
ld -T cm3.ld -o fw crt.o \  
main.o fwlib.a
```

- Converting to the appropriate format

```
objcopy -O binary fw fw.bin
```


Compiling the sources (LLVM)

With LLVM (not cross-compiled):

- Compiling the C source code:

```
clang -emit-llvm -c main.c
```

- Translating the LLVM object:

```
llc -march=thumb -mcpu=cortex-m3 \  
    -o main.s main.ll
```

- Compiling the assembly source:

```
as -o main.o main.s
```

Compiling the sources (LLVM)

- Linking the object files and the driver library (the same as the gcc way):

```
ld -T cm3.ld -o fw crt.o \  
    main.o fwlib.a
```

- Generating binary output file (FLASH image):

```
objcopy -O binary fw fw.bin
```

Behind the scenes: linking

- There are sections in the object files, i.e.:
 - Code (**.code**; the instructions to be run)
 - Data (**.data**, **.bss**; global and static local variables)
 - Read-only data (**.rodata**; const values, initial values of global variables)
 - Debug information (**.debug_info**, **.debug_line**, ...)
 - Interrupt vectors section (**.interrupt_vector** – **defined by us**:

```
__attribute__((section(".interrupt_vector")))
```

Behind the scenes: linking

- Sections in executable files:
 - Sections of the object files should be merged,
 - Sections should be reordered according to the linker script,
 - Function calls, variable access should be resolved,
- Linker scripts
 - Defines the sections (and their order) of the executable files.

Linker scripts

MEMORY

{

FLASH (rx) : ORIGIN = 0x00000000, LENGTH = 0x40000

SRAM (rwx) : ORIGIN = 0x20000000, LENGTH = 0x10000

}

Linker scripts

SECTIONS

{

.text :

{

*(.stack_pointer)

*(.interrupt_vector)

_text = .;

(.text .text.)

(.rodata .rodata)

_etext = .;

} > FLASH

.bss : AT (_etext)

{

_data = .;

(.data .data. .bss .bss.*)

_edata = . ;

} > SRAM

After linking

cortex-m3: file format elf32-littlearm

Sections:

Idx	Name	Size	VMA	LMA	File off	Algn
0	.text	00002674	00000000	00000000	00008000	2**2
		CONTENTS, ALLOC, LOAD, READONLY, CODE				
1	vtable	0000011c	20000000	00002674	00010000	2**2
		CONTENTS, ALLOC, LOAD, DATA				
2	.bss	00000008	2000011c	00002790	0001011c	2**2
		ALLOC, LOAD, DATA				

Debug

- We can use a simulator to debug our code:
 - gdb's built-in simulator (emulates a processor),
 - qemu (emulates a complete device),
- Or we can use the target hardware to debug the firmware; we need
 - a program to download the FLASH content (firmware),
 - a gdb-proxy (which communicates with the hardware),
 - the gdb (GNU's debugger; ARM version),
 - optionally an IDE.

Debugger operations

- Step-by-step C/C++ code analysis
- Breakpoints on instructions
- Watchpoints on variables
- Step-by-step assembly code analysis
- Memory and register dump

Programming languages

Can we use programming languages other than C?

- Industrial computers and system-on-modules boot Linux or Windows CE, thus the answer is: yes.
- C++(0x) code for microcontrollers is not rare, but C# (.NET micro framework) is not really used.
- If the FPGA realizes a soft processor, any languages can be used.

Difficulties with C++

- No memory management
[`operator new()` and `operator delete()` can be implemented],
- Exception handling is problematic,
- Virtual function calls are difficult to implement,
- GNU's `libstdc++` is not enough resource-efficient,
 - But we can use CLANG's `libcxx` instead of GNU's `libstdc++`.

Questions?

Thank you for your attention!

References

- Cortex™-A8 Technical Reference Manual
(Revision: r3p2)
- Cortex™-M3 Technical Reference Manual
(Revision r2p1)
- ARM Architecture Reference Manual ARMv7-A
and ARMv7-R Edition
- ARMv7-M Architecture Reference Manual